

Agent 365 Pilot

Task, Role, Time, Success Plan, Architecture, Next Steps

This plan breaks the Agent 365 pilot into concrete tasks so each team knows what to do, who owns it, roughly how long it takes, what it produces, and how to confirm it worked.

Availability: as of August 28, 2026, all Agent 365 capabilities referenced in this plan are Generally Available (GA); Preview features are excluded.

Sequencing: every task in this plan can be started in parallel, with one exception. Steps that rely on triggering alerts or detections (the Advanced Hunting and custom detection rules in Task 5, and the near-real-time detections in Task 3) need the agents to have generated activity and all the other tasks completed.

How to read this: Each Task groups the steps for one team or workstream, and each step is a card. The title bar names the step, the Control area tag says which control it applies (Governance, Security, Observability, Data security, Compliance, Detection, or Enablement), Target agents says which agents it affects, and the "Question it answers" line states what you are trying to prove. Current state and State after controls applied show the before/after. Then come the team, minimum role(s), estimated hands-on time, the output, and how success is verified. Confirm exact role assignments against your own tenant's model; some tasks can also be done with a broader role such as Global Administrator.

Index

Right-click here and choose Update Field to build the index.

Control & State Summary

Every step summarized by the control it applies, the domain it belongs to, and the before/after state.

Control area	Current state	Domain	Output state
Identity security & governance	Agent identities request tokens with no Conditional Access policy evaluating them, so you cannot see or gate what they access.	Identity & access	A report-only Conditional Access policy evaluates every agent token request and logs a would-be grant/block, giving visibility before enforcement.
Governance (scoped resource access)	Agents reach the resources they use through standing or ambient permissions with no expiry, review, or audit trail.	Identity & access	An Entra ID Governance access package grants each agent scoped, expirable, and auditable access to just the resources it needs.
Governance (identity lifecycle)	If an agent's human sponsor leaves or changes role, the agent can be left ownerless with no automatic reassignment. Review current number of Ownerless Agents in Admin Portal.	Identity & access	A lifecycle workflow automatically notifies the manager and reassigns sponsorship so no agent is left ownerless.
Governance (agent administration)	Before the Agent 365 license is assigned, the AI Administrator has no Agents area in the Microsoft 365 admin center: no Agent Activity view, no approval process, no custom templates, no risky agents signal, and no way to block an agent.	Governance	With the Agent 365 license, the AI Administrator gets the Agents area with Agent Activity (usage and telemetry), the approval process, custom templates, the risky agents view, and the ability to block agents.
Governance (approval & policy attachment)	Agents are pending approval, guardrails are not yet attached, and the wider org cannot install them.	Governance	Both agents are approved and published with the CA policy and access package attached at activation.
Security (visibility / posture)	The SOC has no single inventory or posture view of managed agents in Defender.	Security	Both pilot agents appear in the Defender AI agent inventory with posture details and Agent 365 shows Connected.
Security (runtime protection)	Risky agent tool invocations run without runtime evaluation or the ability to block them.	Security	Defender can evaluate agent-initiated tool invocations and block risky actions before they execute.
Observability (log ingestion) + Security	Agent activity is not flowing into the hunting tables, so alerts cannot fire and the SOC cannot query it.	Security	Agent activity flows into CloudAppEvents and AgentsInfo, and near-real-time alerts begin firing on suspicious behavior.
Data security (classification)	Content the agents consume is unclassified, so it cannot be selectively governed.	Data security & compliance	Confidential and General labels are published and selectable, enabling classification-based controls.
Data security (DLP / oversharing)	Agents can surface Confidential content directly in their answers.	Data security & compliance	Agents withhold Confidential values in answers and each match is logged in Purview DLP alerts.
Data security (DLP / exfiltration)	Agents can email financial PII (credit card / ABA routing numbers) out of the org.	Data security & compliance	Outbound agent mail carrying financial PII is blocked before delivery and each attempt is logged.
Compliance / Insider risk	Unethical or risky agent interactions are not detected or routed to a reviewer.	Data security & compliance	Unethical interactions are flagged to a reviewer queue and risky agent activity is scored by Insider Risk.
Observability (inventory)	There is no queryable, owner-resolved inventory of the agents.	Detection & response	A KQL inventory resolves each agent to its owner, platform, model, and lifecycle status.
Observability (activity)	Agent actions across Microsoft 365 are not visible in one place.	Detection & response	Queries show each agent's install, usage, and tool-invocation events with initiator and client IP.
Threat detection (proactive hunting)	A compromised or misconfigured agent can go unnoticed with no proactive hunt.	Detection & response	Hunting queries surface failed-sign-in spikes, agents missing guardrails, and orphaned agents for triage.

Control area	Current state	Domain	Output state
Detection (automation / alerting)	The SOC's best hunts run manually, if at all, with no scheduled alerting.	Detection & response	Saved hunts run on a schedule as custom detection rules that raise alerts and incidents automatically.
Governance (third-party onboarding)	External Genie agents sit outside the Agent 365 registry and its governance.	Ecosystem / third-party	A validated Registry sync connection links the Databricks workspace to the Agent 365 registry.
Enablement (governed end-user access)	End users cannot chat with the governed Genie agents from inside Teams.	Ecosystem / third-party	A governed Genie app is available in Teams and answers data questions with its sources shown.
Governance (third-party onboarding)	External Vertex agents sit outside the Agent 365 registry.	Ecosystem / third-party	A validated Connected platforms entry links the Vertex AI project to the Agent 365 registry.

Estimated Time Rollup

Task	Estimated hands-on time
Task 1 - Agent Identities & Guardrails	~2 hours
Task 2 - Governance & AI Admin controls	~50 minutes
Task 3 - Connect to Security Portal	~1.5 - 2 hours
Task 4 - Secure Data (Purview)	~1.75 hours
Task 5 - Advanced Hunting	~2 hours
Task 6 - Connected Platforms (Databricks Genie; optional Marketplace + Google Vertex)	~20 min core (+20 min Marketplace, +40 min Vertex)
Total core (excludes optional items, ingestion/propagation waits, and roadmap items)	~8.5 - 9 hours core (~9.5 with optional Marketplace + Vertex)

Task 1 - Agent Identities & Guardrails (Identity team)

The Identity Admin applies guardrails to the agent identities in Microsoft Entra. These guardrails are later consumed by the AI Admin at publish time (Task 2).

Step 1.1 Apply Conditional Access policy (Report-only) to the Copilot Studio and AI Foundry agent identities	
Control area: Identity security & governance	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)
Question it answers	<i>Can we see, and later control, every token our agents request before we turn on enforcement?</i>
Current state	Agent identities request tokens with no Conditional Access policy evaluating them, so you cannot see or gate what they access.
State after controls applied	A report-only Conditional Access policy evaluates every agent token request and logs a would-be grant/block, giving visibility before enforcement.
Team	Identity
Role(s) required	Conditional Access Administrator (or Security Administrator)
Est. time	45 min
Output / deliverable	Create CA policy i.e. "Observe - Pilot agents access" scoped to both pilot agent identities (or their blueprint principals), Report-only.
What success looks like	Agents keep working. Entra ID > Monitoring > Sign-in logs > Service principal sign-ins shows the policy evaluated with a Report-only "would have been granted/blocked" result for each agent token request. No agent disruption.

Step 1.2 Provision an Entra ID Governance access package for agent resource access	
Control area: Governance (scoped resource access)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)
Question it answers	<i>Do our agents get only the scoped, time-bound, auditable access they need, instead of standing permissions?</i>
Current state	Agents reach the resources they use through standing or ambient permissions with no expiry, review, or audit trail.
State after controls applied	An Entra ID Governance access package grants each agent scoped, expirable, and auditable access to just the resources it needs.
Team	Identity
Role(s) required	Identity Governance Administrator
Est. time	30 min
Output / deliverable	Create an Access package i.e. "Agent365 Pilot - agent resource access" (in a dedicated catalog) with the required resource roles, assigned to each agent identity with an expiration and access review.
What success looks like	Each agent's access is granted through the package (Identity Governance > Entitlement management > Access packages), is time-bound and reviewable, and can be revoked from one place.

Step 1.3 Configure sponsorship-change Lifecycle Workflow (ownerless-agent reassignment)	
Control area: Governance (identity lifecycle)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)
Question it answers	<i>If an agent's human owner leaves or changes role, does the agent automatically get a new accountable owner?</i>
Current state	If an agent's human sponsor leaves or changes role, the agent can be left ownerless with no automatic reassignment. Review current number of Ownerless Agents in Admin Portal.
State after controls applied	A lifecycle workflow automatically notifies the manager and reassigns sponsorship so no agent is left ownerless.
Team	Identity
Role(s) required	Lifecycle Workflows Administrator (ID Governance)
Est. time	45 min
Output / deliverable	Workflow "Notify manager - agent sponsorship change" built from the agent-sponsor template, with tasks: email manager + transfer sponsorship.
What success looks like	Workflow history > Tasks shows both tasks Successful; the manager receives the sponsorship-change email; the agent's Sponsor is reassigned automatically so no agent is left ownerless.

Task 1 subtotal: ~2 hours

Task 2 - Apply Governance at Publish / Approval (AI Admin)

Acting as the AI Admin, use the controls the Agent 365 license unlocks (Agent Activity, the approval process, custom templates, risky agents, and blocking agents), then run the approval workflow and attach the Task 1 guardrails at publish so agents are compliant before the wider org consumes them.

Step 2.1 Use the AI Administrator controls unlocked by the Agent 365 license	
Control area: Governance (agent administration)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)

Question it answers	<i>Once the Agent 365 license is assigned, can the AI Administrator see and govern agents from the admin center?</i>
Current state	Before the Agent 365 license is assigned, the AI Administrator has no Agents area in the Microsoft 365 admin center: no Agent Activity view, no approval process, no custom templates, no risky agents signal, and no way to block an agent.
State after controls applied	With the Agent 365 license, the AI Administrator gets the Agents area with Agent Activity (usage and telemetry), the approval process, custom templates, the risky agents view, and the ability to block agents.
Team	AI Admin
Role(s) required	AI Administrator
Est. time	20 min
Output / deliverable	AI Administrator walkthrough of the Agents area: Agent Activity reviewed, approval process located, at least one custom template created, risky agents view checked, and the block-agent action confirmed.
What success looks like	The AI Administrator can open Agent Activity, run the approval process, create a custom template, see the risky agents list, and block an agent, all from the Microsoft 365 admin center.

Step 2.2 Approve pending agents and apply governance at publish (CA policy + access package)	
Control area: Governance (approval & policy attachment)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)
Question it answers	<i>Are our agents compliant with our policies before the rest of the org can install and use them?</i>
Current state	Agents are pending approval, guardrails are not yet attached, and the wider org cannot install them.
State after controls applied	Both agents are approved and published with the CA policy and access package attached at activation.
Team	AI Admin
Role(s) required	AI Administrator
Est. time	30 min
Output / deliverable	Both pilot agents published to the org store via the Default template, with the "Observe - Pilot agents access" CA policy and the agent access package attached at activation.
What success looks like	Entra Enterprise app for each agent shows the assigned access package; agent sign-ins show the CA policy applied (Report-only). Governance is attached because policies apply to new activations only, so this is the one moment to add them.

Task 2 subtotal: ~50 minutes

Task 3 - Connect Agent 365 to the Security Portal (SOC / Security)

Connect Agent 365 to Microsoft Defender so the SOC gets one agent inventory, near-real-time alerts, runtime blocking of unsafe tool invocations, and hunting data.

Step 3.1 Enable Security for AI agents in Defender	
Control area: Security (visibility / posture)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)
Question it answers	<i>Can our SOC see every managed agent and its security posture in one inventory?</i>
Current state	The SOC has no single inventory or posture view of managed agents in Defender.

State after controls applied	Both pilot agents appear in the Defender AI agent inventory with posture details and Agent 365 shows Connected.
Team	SOC / Security
Role(s) required	Security Administrator
Est. time	20 min
Output / deliverable	Security for AI agents toggled on; Agent 365 shows Connected under AI real-time protection & investigation.
What success looks like	Both pilot agents appear in Assets, Agents, inventory with posture details; Agent 365 shows Connected, sample alerts will be provided.

Step 3.2 Enable extended detections for Copilot Studio Agents	
Control area: Security (runtime protection)	
Target agents	Copilot Studio - Power Platform Admin Center.
Question it answers	<i>Can we detect and block risky agent tool invocations before they actually run?</i>
Current state	Risky agent tool invocations run without runtime evaluation or the ability to block them.
State after controls applied	Defender can evaluate agent-initiated tool invocations and block risky actions before they execute.
Team	SOC / Security
Role(s) required	Security Administrator + Power Platform Administrator (to register the Copilot webhook Entra app)
Est. time	45-60 min
Output / deliverable	Threat protection enabled; Copilot Studio real-time protection connected via a registered Entra app (App ID + Federated Identity Credential) and the Power Platform integration endpoint.
What success looks like	Real-time protection shows On; Defender can evaluate agent-initiated tool invocations and block risky actions before they execute.

Step 3.3 Connect the Microsoft 365 app connector for near-real-time detections	
Control area: Observability (log ingestion) + Security	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)
Question it answers	<i>Is agent activity flowing into our hunting tables so alerts can fire and the SOC can query it?</i>
Current state	Agent activity is not flowing into the hunting tables, so alerts cannot fire and the SOC cannot query it.
State after controls applied	Agent activity flows into CloudAppEvents and AgentsInfo, and near-real-time alerts begin firing on suspicious behavior.
Team	SOC / Security
Role(s) required	Application Administrator or Cloud Application Administrator (Purview auditing turned on)
Est. time	30 min
Output / deliverable	Microsoft 365 connector shows Connected in Defender for Cloud Apps, with all components selected.
What success looks like	Agent activity flows into the CloudAppEvents and AgentsInfo Advanced Hunting tables; near-real-time alerts begin to fire on suspicious agent behavior.

Task 3 subtotal: ~1.5 - 2 hours (plus first-time ingestion wait)

Task 4 - Secure the Data Agents Consume (Compliance team)

Wearing the Compliance hat, apply Microsoft Purview guardrails so the agents cannot overshare or leak sensitive data. (Data-source setup such as building the SharePoint site and grounding the agent is a prerequisite handled with the agent owner.)

Step 4.1 Publish sensitivity labels (Confidential, General)	
Control area: Data security (classification)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent) (and the content they ground on)
Question it answers	<i>Can we classify which content agents may or may not use?</i>
Current state	Content the agents consume is unclassified, so it cannot be selectively governed.
State after controls applied	Confidential and General labels are published and selectable, enabling classification-based controls.
Team	Compliance
Role(s) required	Compliance Administrator
Est. time	30 min
Output / deliverable	Two published labels in a label policy scoped to files and email.
What success looks like	Both labels appear under Label policies and are selectable from the Sensitivity ribbon in Word/Excel/SharePoint.

Step 4.2 DLP #1 - block Copilot / agents from processing Confidential content	
Control area: Data security (DLP / oversharing)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)
Question it answers	<i>Can agents be stopped from surfacing Confidential content in their answers?</i>
Current state	Agents can surface Confidential content directly in their answers.
State after controls applied	Agents withhold Confidential values in answers and each match is logged in Purview DLP alerts.
Team	Compliance
Role(s) required	Compliance Administrator
Est. time	30 min
Output / deliverable	DLP policy "Block Copilot on Confidential" scoped to M365 Copilot + Copilot Chat, restricting access to knowledge sources labeled Confidential.
What success looks like	Agent withholds confidential values in answers (file may still be cited, but no names/amounts/card numbers returned) and Purview > Data Loss Prevention > Alerts logs a matching "Deny grounding on Confidential" event.

Step 4.3 DLP #2 - block email of financial PII (credit card / ABA routing numbers)	
Control area: Data security (DLP / exfiltration)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)
Question it answers	<i>Can agents be stopped from emailing financial PII out of the org?</i>
Current state	Agents can email financial PII (credit card / ABA routing numbers) out of the org.
State after controls applied	Outbound agent mail carrying financial PII is blocked before delivery and each attempt is logged.
Team	Compliance
Role(s) required	Compliance Administrator
Est. time	30 min
Output / deliverable	DLP policy "Block email of financial PII" scoped to Exchange, blocking outbound mail carrying Credit Card or ABA Routing numbers.

What success looks like	The agent's Outlook Send-email action fails, the mail is never delivered, and Purview > DLP > Alerts logs a "DLP policy match for email..." event with the financial sensitive info types.
--------------------------------	--

Step 4.4 Turn on Communication Compliance and confirm Insider Risk default agent policy	
Control area: Compliance / Insider risk	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)
Question it answers	<i>Are unethical or risky agent interactions detected and put in front of a reviewer?</i>
Current state	Unethical or risky agent interactions are not detected or routed to a reviewer.
State after controls applied	Unethical interactions are flagged to a reviewer queue and risky agent activity is scored by Insider Risk.
Team	Compliance
Role(s) required	Compliance Administrator
Est. time	20 min
Output / deliverable	Communication Compliance policy "Detect unethical interactions for AI agents" scoped to Copilot Studio and Azure Foundry; default Insider Risk Management agent policy confirmed enabled.
What success looks like	Unethical prompts/responses are flagged to the reviewer queue within minutes; risky agent activity is scored by Insider Risk with no extra configuration.

Task 4 subtotal: ~1.75 hours (source-grounding prerequisites add time)

Task 5 - Analyze Agent Logs in Advanced Hunting (SOC)

Acting as the SOC, query and analyze the raw agent logs with KQL in the Defender portal, then promote the best queries to custom detection rules. Requires Task 3 complete and both agents exercised so there is real activity.

Step 5.1 Inventory agents with KQL (AgentsInfo table)	
Control area: Observability (inventory)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent) (all managed agents)
Question it answers	<i>Do we have an accurate inventory of our agents and who owns them?</i>
Current state	There is no queryable, owner-resolved inventory of the agents.
State after controls applied	A KQL inventory resolves each agent to its owner, platform, model, and lifecycle status.
Team	SOC
Role(s) required	Security Reader
Est. time	20 min
Output / deliverable	Inventory query returning each agent with owner (resolved via IdentityInfo), platform, model, published/lifecycle status.
What success looks like	Both pilot agents appear resolved to the expected owner UPNs; no unknown agents and no ownerless (all-zero owner) agents unaccounted for.

Step 5.2 Review agent activity and tool invocations	
Control area: Observability (activity)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent)
Question it answers	<i>Can we see what each agent is actually doing across Microsoft 365?</i>
Current state	Agent actions across Microsoft 365 are not visible in one place.

State after controls applied	Queries show each agent's install, usage, and tool-invocation events with initiator and client IP.
Team	SOC
Role(s) required	Security Reader
Est. time	20 min
Output / deliverable	Queries over CloudAppEvents / OfficeActivity / MessageEvents showing Teams app/bot install and usage events per agent.
What success looks like	App/bot install and membership events line up with the pilot rollout; every initiator (UserId) and ClientIP is expected.

Step 5.3 Hunt for suspicious agent behavior (brute-force sign-ins, missing guardrails, orphaned agents)	
Control area: Threat detection (proactive hunting)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent) (all managed agents)
Question it answers	<i>Can we proactively find a compromised or misconfigured agent?</i>
Current state	A compromised or misconfigured agent can go unnoticed with no proactive hunt.
State after controls applied	Hunting queries surface failed-sign-in spikes, agents missing guardrails, and orphaned agents for triage.
Team	SOC
Role(s) required	Security Reader
Est. time	45 min
Output / deliverable	Three hunting queries: failed sign-in spikes (SigninLogs), agents missing Instructions/Guardrails (AgentsInfo), and unowned-but-deployed agents (AgentsInfo).
What success looks like	Each hunt returns actionable rows or a clean, explainable empty set; misconfigured/orphaned agents surface with the fields needed to triage.

Step 5.4 Save queries and create custom detection rules	
Control area: Detection (automation / alerting)	
Target agents	Both pilot agents (Copilot Studio agent + AI Foundry agent) (all managed agents)
Question it answers	<i>Are the SOC's best hunts turned into automatic alerts instead of manual checks?</i>
Current state	The SOC's best hunts run manually, if at all, with no scheduled alerting.
State after controls applied	Saved hunts run on a schedule as custom detection rules that raise alerts and incidents automatically.
Team	SOC
Role(s) required	Security Operator or Security Administrator
Est. time	30 min
Output / deliverable	Saved hunting queries plus at least one custom detection rule (frequency, severity, impacted entities mapped).
What success looks like	Defender runs the rule on schedule and raises alerts/incidents the SOC can triage automatically.

Task 5 subtotal: ~2 hours

Task 6 - Connected Platforms (AI Admin + platform admins)

Connect external agent platforms to Agent 365 through Registry sync (the Connected platforms feature in the Microsoft 365 admin center) so their agents show up in the Agent 365 registry alongside the pilot agents. Databricks Genie is part of the pilot; Google Vertex AI is optional for the customer.

Step 6.1 Connect Databricks Genie in Registry sync	
Control area: Governance (third-party onboarding)	
Target agents	Databricks Genie agents
Question it answers	<i>Are our external Genie agents connected to the Agent 365 agent registry?</i>
Current state	External Genie agents sit outside the Agent 365 registry and its governance.
State after controls applied	A validated Registry sync connection links the Databricks workspace to the Agent 365 registry.
Team	AI Admin
Role(s) required	AI Administrator
Est. time	20 min
Output / deliverable	A validated, saved Registry sync connection (Workspace URL + Client ID + Secret + region) in the M365 admin center.
What success looks like	Validate succeeds (Agent 365 can reach and authenticate to the workspace) and the connection saves without error.

Step 6.2 (Optional) Install Databricks Genie from the AI Marketplace and configure it in Teams	
Control area: Enablement (governed end-user access)	
Target agents	Databricks Genie agents (end users in Teams)
Question it answers	<i>Can end users chat with governed Genie agents directly in Teams?</i>
Current state	End users cannot chat with the governed Genie agents from inside Teams.
State after controls applied	A governed Genie app is available in Teams and answers data questions with its sources shown.
Team	AI Admin
Role(s) required	AI Administrator
Est. time	20 min
Output / deliverable	Marketplace app made available to users; a Teams chat connected to a specific workspace + Genie space.
What success looks like	The app shows Available in Agents > All agents; a configured Teams chat answers data questions and shows Sources. Genie activities will be available to the SOC team in Advanced Hunting.

Step 6.3 (Optional) Connect Google Vertex AI in Registry sync	
Control area: Governance (third-party onboarding)	
Target agents	Google Vertex AI agents (Agent Engine / Reasoning Engines)
Question it answers	<i>Are our external Vertex agents connected to the Agent 365 agent registry?</i>
Current state	External Vertex agents sit outside the Agent 365 registry.
State after controls applied	A validated Connected platforms entry links the Vertex AI project to the Agent 365 registry.
Team	AI Admin
Role(s) required	AI Administrator + Google Cloud Admin
Est. time	40 min
Output / deliverable	First Sync agents completed on the Vertex connection; imported Vertex agents visible in the registry.
What success looks like	Connection shows Last sync status = success; Vertex agents appear in the registry alongside the Copilot Studio, Foundry, and Genie agents.

Task 6 subtotal: ~20 minutes (Databricks Genie connect); optional Marketplace install +20 min, optional Google Vertex +40 min

Final State for the Pilot

When every task above is complete, the pilot has reached the following end state. This is the success definition for the pilot as a whole, across all teams.

- Every pilot agent has a governed Microsoft Entra identity with a Conditional Access policy evaluating its tokens, a scoped and expirable access package for the resources it uses, and an automatic sponsorship-reassignment workflow so no agent is ever left ownerless.
- The AI Administrator uses the controls the Agent 365 license unlocks - Agent Activity, approvals, custom templates, risky agents, and blocking agents - and agents were approved and published with the guardrails attached at activation, so they were compliant before the wider org could install and use them.
- The SOC sees every managed agent and its posture in one Microsoft Defender inventory, with runtime protection that can block risky tool invocations before they execute and agent activity flowing into the hunting tables.
- The data agents consume is protected by Microsoft Purview: sensitivity labels classify content, DLP stops agents from surfacing Confidential content and from emailing financial PII, and unethical interactions are flagged to reviewers and scored by Insider Risk.
- The SOC has an owner-resolved agent inventory, visibility into agent activity, proactive hunts for compromised or misconfigured agents, and the best hunts promoted to scheduled custom detection rules that raise alerts automatically.
- External agent platforms are onboarded through Registry sync's Connected platforms - Databricks Genie in the pilot, and optionally Google Vertex AI - and appear in the same Agent 365 registry as the pilot agents, so the organization has one central, governed inventory of first-party and third-party agents.

Roadmap

Roadmap sessions to be included during the Pilot.

Training

Training sessions to be included during the Pilot.

Roles & Teams at a Glance

Verify these against your tenant's role model; some tasks can also be performed with Global Administrator.

Team	Typical roles used in this pilot	Tasks
------	----------------------------------	-------

Team	Typical roles used in this pilot	Tasks
Identity team	Conditional Access Administrator, Identity Governance Administrator, Lifecycle Workflows Administrator	Task 1
AI Admin	AI Administrator (agent approval, agent controls, Connected platforms / registry sync)	Tasks 2, 6
SOC / Security	Security Administrator, Security Operator, Security Reader, Application/Cloud Application Administrator, Power Platform Administrator	Tasks 3, 5, roadmap
Compliance team	Compliance Administrator	Task 4
Google Cloud / Platform admin (optional)	Google Cloud IAM admin (project owner or equivalent)	Task 6 (optional Vertex)

Roles Summary

Every step in this plan with its owning role, team, and estimated hands-on time.

Role	Task	Team	Est. time
Conditional Access Administrator (or Security Administrator)	1.1 Apply Conditional Access policy (Report-only) to the Copilot Studio and AI Foundry agent identities	Identity	45 min
Identity Governance Administrator	1.2 Provision an Entra ID Governance access package for agent resource access	Identity	30 min
Lifecycle Workflows Administrator (ID Governance)	1.3 Configure sponsorship-change Lifecycle Workflow (ownerless-agent reassignment)	Identity	45 min
AI Administrator	2.1 Use the AI Administrator controls unlocked by the Agent 365 license	AI Admin	20 min
AI Administrator	2.2 Approve pending agents and apply governance at publish (CA policy + access package)	AI Admin	30 min
Security Administrator	3.1 Enable Security for AI agents in Defender	SOC / Security	20 min
Security Administrator + Power Platform Administrator (to register the Copilot webhook Entra app)	3.2 Enable extended detections for Copilot Studio Agents	SOC / Security	45-60 min
Application Administrator or Cloud Application Administrator (Purview auditing turned on)	3.3 Connect the Microsoft 365 app connector for near-real-time detections	SOC / Security	30 min
Compliance Administrator	4.1 Publish sensitivity labels (Confidential, General)	Compliance	30 min
Compliance Administrator	4.2 DLP #1 - block Copilot / agents from processing Confidential content	Compliance	30 min
Compliance Administrator	4.3 DLP #2 - block email of financial PII (credit card / ABA routing numbers)	Compliance	30 min
Compliance Administrator	4.4 Turn on Communication Compliance and confirm Insider Risk default agent policy	Compliance	20 min
Security Reader	5.1 Inventory agents with KQL (AgentsInfo table)	SOC	20 min
Security Reader	5.2 Review agent activity and tool invocations	SOC	20 min
Security Reader	5.3 Hunt for suspicious agent behavior (brute-force sign-ins, missing guardrails, orphaned agents)	SOC	45 min
Security Operator or Security Administrator	5.4 Save queries and create custom detection rules	SOC	30 min
AI Administrator	6.1 Connect Databricks Genie in Registry sync	AI Admin	20 min
AI Administrator	6.2 (Optional) Install Databricks Genie from the AI Marketplace and configure it in Teams	AI Admin	20 min
AI Administrator + Google Cloud Admin	6.3 (Optional) Connect Google Vertex AI in Registry sync	AI Admin	40 min

Architecture Diagram

Out of Scope

The following are out of scope, and can be included in Phase I of Deployment

- Shadow AI, discovering/blocking local agents in endpoints (Claude Desktop, ChatGPT, OpenClaw), discovering code agents like Claude code, cursor or others.
- Entra Suite, GSA, analyzing Network traffic, Agents.
- Production enforcement and org-wide rollout - moving the Conditional Access policy from Report-only to Enforce and rolling governance out beyond the pilot scope.
- Agent 365 SDK for third party agents.

Next Steps, post Pilot

- Licensing